


CS183

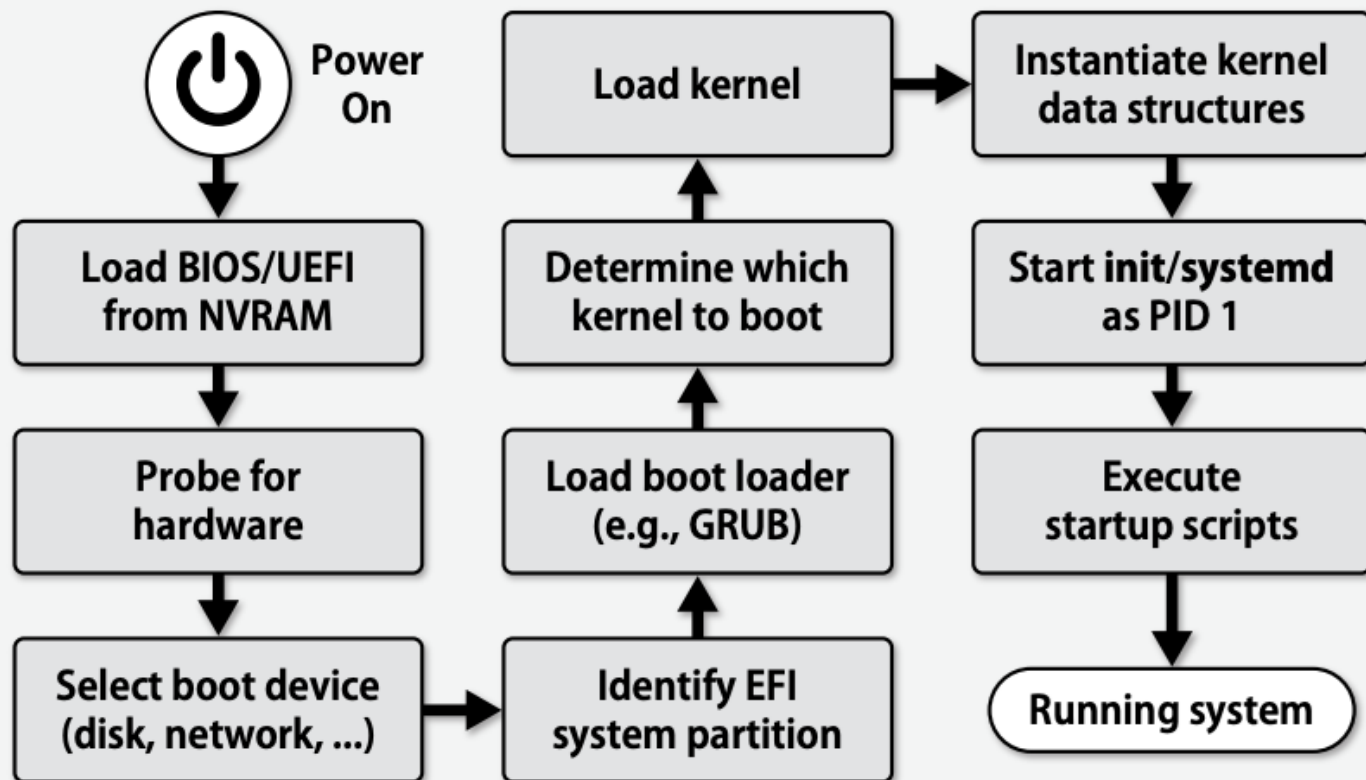
Instructor: Ali Davanian

(Slides were adopted from Brian Crites and Alireza Abdoli)

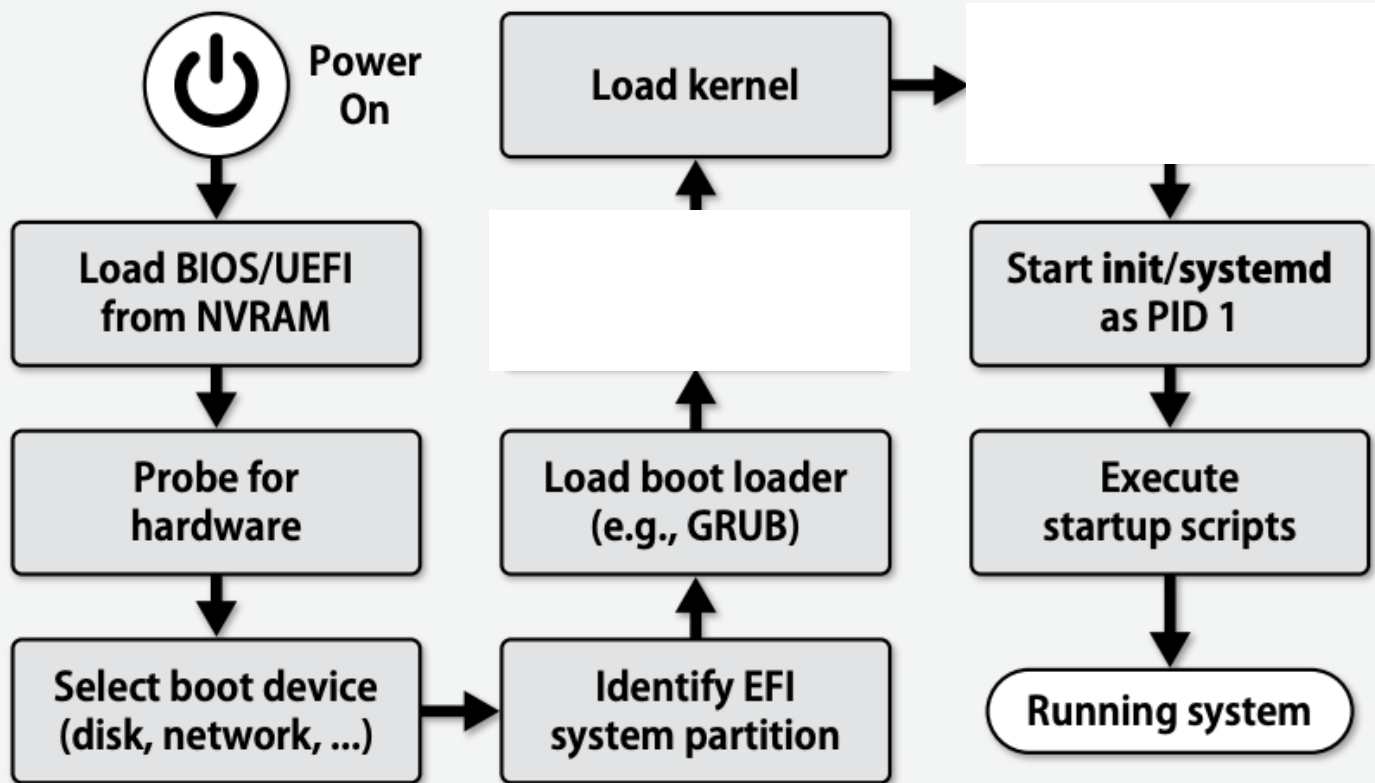


Linux Boot Systems

Linux & UNIX boot process

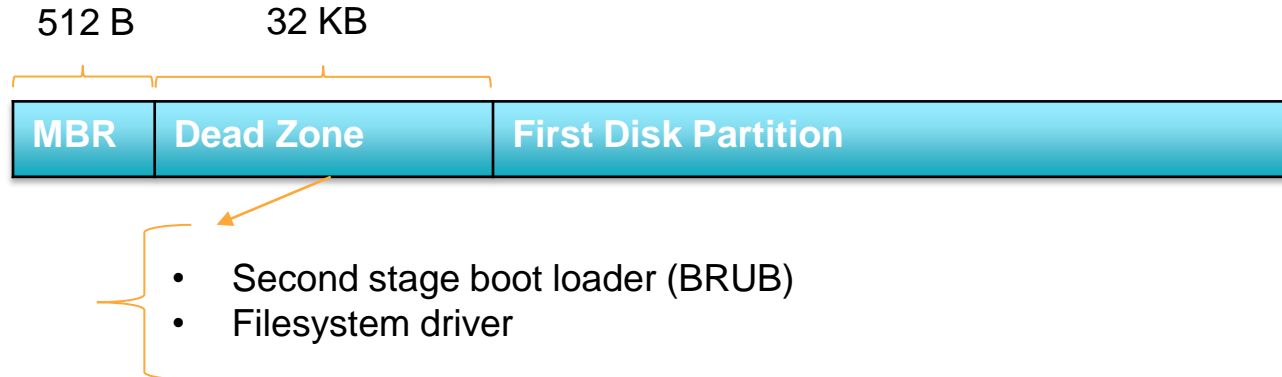


Linux & UNIX boot process



Three Booting Stages (Basic I/O System - BIOS)

- (1) Master Boot Record (MBR): Points to the secondary boot loader
- (2) Secondary Boot Loader
 - (a) Usually GRand Unified Bootloader (GRUB) which lets you choose which partition to load as the OS
 - (b) GRUB loads the secondary bootloader for the specific OS kernel
- (3) OS Kernel: The core functionality for your operating system



Unified Extensible Firmware Interface (UEFI)

- BIOS is being replaced by Unified Extensible Firmware Interface (**UEFI**) which covers both MBR and secondary boot loader with some extra features:
 - It understands FAT file system
 - Less bootstrapping
 - UEFI has a formal API that allows accessing the system's hardware and custom configurations

GRUB

- Understands **file systems** giving it access to enough memory to load the OS
- Reads a configuration file usually at `/boot/grub(2)/grub.cfg`
- Can be generated with the utility `grub(2)-mkconfig`
- Config is often re-created during OS updates, so you may need to take action to stop custom `grub.cfg` files from being overwritten

Common GRUB configuration options from `/etc/default/grub`

Shell variable name	Contents or function
GRUB_BACKGROUND	Background image ^a
GRUB_CMDLINE_LINUX	Kernel parameters to add to menu entries for Linux ^b
GRUB_DEFAULT	Number or title of the default menu entry
GRUB_DISABLE_RECOVERY	Prevents the generation of recovery mode entries
GRUB_PRELOAD_MODULES	List of GRUB modules to be loaded as early as possible
GRUB_TIMEOUT	Seconds to display the boot menu before autoboot

a. The background image must be a `.png`, `.tga`, `.jpg`, or `.jpeg` file.

b. Table 2.3 on page 38 lists some of the available options.

Kernel Booting (Init/Systemd)

- Kernel starts a single init or systemd process with a PID of 1
- This process spawns all other processes in the system (using fork)
- Options for booting specified in `/etc/grub.d/40_custom`

Examples of kernel boot time options

Option	Meaning
debug	Turns on kernel debugging
init=/bin/bash	Starts only the bash shell; useful for emergency recovery
root=/dev/foo	Tells the kernel to use /dev/foo as the root device
single	Boots to single-user mode

```
$ ps -aux
```

```
USER      PID %CPU %MEM    USZ    RSS TTY      STATE START   TIME COMMAND
root         1  0.0  0.3 125488  3876 ?        Ss   10:37   0:01 /usr/lib/systemd/systemd --switched
--root --system --deserialize 22
root         2  0.0  0.0      0      0 ?        S    10:37   0:00 [kthreadd]
```

Kernel has many components known as “UNIT”s

```
UNIT                                LOAD    ACTIVE SUB    DESCRIPTION
proc-sys-fs-binfmt_misc.automount loaded active running Arbitrary Executable File Formats File
sys-devices-pci0000:00-0000:00:01.1-ata2-host1-target1:0:0-1:0:0-0-block-sr0.device loaded active
sys-devices-pci0000:00-0000:00:03.0-net-ens3.device loaded active plugged 82540EM Gigabit Ethernet
sys-devices-pci0000:00-0000:00:05.0-sound-card0.device loaded active plugged 82801AA AC'97 Audio
sys-devices-pci0000:00-0000:00:0d.0-ata3-host2-target2:0:0-2:0:0-0-block-sda-sda1.device loaded active
sys-devices-pci0000:00-0000:00:0d.0-ata3-host2-target2:0:0-2:0:0-0-block-sda-sda2.device loaded active
sys-devices-pci0000:00-0000:00:0d.0-ata3-host2-target2:0:0-2:0:0-0-block-sda-sda3.device loaded active
sys-devices-platform-serial8250-tty-ttyS0.device loaded active plugged /sys/devices/platform/serial8250/tty/ttyS0
sys-devices-platform-serial8250-tty-ttyS1.device loaded active plugged /sys/devices/platform/serial8250/tty/ttyS1
sys-devices-platform-serial8250-tty-ttyS2.device loaded active plugged /sys/devices/platform/serial8250/tty/ttyS2
sys-devices-platform-serial8250-tty-ttyS3.device loaded active plugged /sys/devices/platform/serial8250/tty/ttyS3
sys-devices-virtual-block-dm0-dm0.device loaded active plugged /sys/devices/virtual/block/dm0
sys-devices-virtual-block-dm1-dm1.device loaded active plugged /sys/devices/virtual/block/dm1
sys-module-configfs.device loaded active plugged /sys/module/configfs
sys-module-fuse.device loaded active plugged /sys/module/fuse
sys-subsystem-net-devices-ens3.device loaded active plugged 82540EM Gigabit Ethernet Controller
-.mount loaded active mounted /
boot.mount loaded active mounted /boot
dev-hugepages.mount loaded active mounted Huge Pages File System
dev-mqueue.mount loaded active mounted POSIX Message Queue File System
proc-sys-fs-binfmt_misc.mount loaded active mounted Arbitrary Executable File Formats File
run-user-1000.mount loaded active mounted /run/user/1000
sys-kernel-config.mount loaded active mounted Configuration File System
sys-kernel-debug.mount loaded active mounted Debug File System
systemd-ask-password-plymouth.path loaded active waiting Forward Password Requests to Plymouth
systemd-ask-password-wall.path loaded active waiting Forward Password Requests to Wall Directory Watch
session-3.scope loaded active running Session 3 of user sina
auditd.service loaded active running Security Auditing Service
crond.service loaded active running Command Scheduler
dbus.service loaded active running D-Bus System Message Bus
firewalld.service loaded active running firewalld - dynamic firewall daemon
getty@tty1.service loaded active running Getty on tty1
kdump.service loaded failed failed Crash recovery kernel arming
kmod-static-nodes.service loaded active exited Create list of required static device nodes
lvm2-lvmetad.service loaded active running LVM2 metadata daemon
```

lines 1-36

- Services
- Sockets
- Devices
- Mount points
- Partitions
- Timers
- etc.

Systemd

- Anything managed by systemd is known as a “unit”
- Each unit describes some boot item or a boot requirement
 - services, sockets, devices, mount points, partitions, timers, etc.
- systemd units are located at:
 - /lib/systemd/system
 - /usr/lib/systemd/system
 - /etc/systemd/system

```
[Unit]
```

```
Description=fast remote file copy program daemon
```

```
ConditionPathExists=/etc/rsyncd.conf
```

```
[Service]
```

```
ExecStart=/usr/bin/rsync --daemon --no-detach
```

```
[Install]
```

```
WantedBy=multi-user.target
```

Unit Files

- Describes dependencies between other units and targets
- Describes relative ordering between dependent units, allowing for parallelism
- Describes various modes of operation
- Can handle preparation and post-death cleanup
- Lots more...

```
[Unit]
Description=The nginx HTTP and reverse proxy server
After=network.target remote-fs.target nss-lookup.target

[Service]
Type=forking
PIDFile=/run/nginx.pid
ExecStartPre=/usr/bin/rm -f /run/nginx.pid
ExecStartPre=/usr/sbin/nginx -t
ExecStart=/usr/sbin/nginx
ExecReload=/bin/kill -s HUP $MAINPID
KillMode=process
KillSignal=SIGQUIT
TimeoutStopSec=5
PrivateTmp=true

[Install]
WantedBy=multi-user.target
```

Unit Files

[Unit]

Description=The nginx HTTP and reverse proxy server

After=network.target remote-fs.target nss-lookup.target

[Service]

Type=forking

PIDFile=/run/nginx.pid

ExecStartPre=/usr/bin/rm -f /run/nginx.pid

ExecStartPre=/usr/sbin/nginx -t

ExecStart=/usr/sbin/nginx

ExecReload=/bin/kill -s HUP \$MAINPID

KillMode=process

KillSignal=SIGQUIT

TimeoutStopSec=5

PrivateTmp=true

[Install]

WantedBy=multi-user.target

- [Unit] Section

- Starts with a description of the process
- After lists other targets which this process requires to be able to run

Unit Files

[Unit]

Description=The nginx HTTP and reverse proxy server

After=network.target remote-fs.target nss-lookup.target

[Service]

Type=forking

PIDFile=/run/nginx.pid

ExecStartPre=/usr/bin/rm -f /run/nginx.pid

ExecStartPre=/usr/sbin/nginx -t

ExecStart=/usr/sbin/nginx

ExecReload=/bin/kill -s HUP \$MAINPID

KillMode=process

KillSignal=SIGQUIT

TimeoutStopSec=5

PrivateTmp=true

[Install]

WantedBy=multi-user.target

- [Service] Section

- Type defines its running condition, here forking means the daemon should keep running even though the process finishes
- Daemon PID is recorded in PIDFile location, since forking means systemd won't have created the daemon directly
- Exec lines show commands to run during various forms of execution (pre, start, reload, etc.)
- KillMode and KillSignal state that a QUIT signal should terminate and clean up
- PrivateTmp is for security reasons and puts the /tmp directory someplace users can't freely access it

Unit Files

[Unit]

Description=The nginx HTTP and reverse proxy server

After=network.target remote-fs.target nss-lookup.target

[Service]

Type=forking

PIDFile=/run/nginx.pid

ExecStartPre=/usr/bin/rm -f /run/nginx.pid

ExecStartPre=/usr/sbin/nginx -t

ExecStart=/usr/sbin/nginx

ExecReload=/bin/kill -s HUP \$MAINPID

KillMode=process

KillSignal=SIGQUIT

TimeoutStopSec=5

PrivateTmp=true

[Install]

WantedBy=multi-user.target

- [Install] Section

- WantedBy links this unit with a target, so when the target is enabled this unit will be linked
- This allows the target to stay relatively clean but to be extended from the outside

systemctl

- Allows you to investigate and modify systemd
- Systemd unit file describes what it does, systemctl enables/disable it to run at boot

Commonly used systemctl subcommands

Subcommand	Function
<code>list-unit-files</code> [<i>pattern</i>]	Shows installed units; optionally matching <i>pattern</i>
<code>enable</code> <i>unit</i>	Enables <i>unit</i> to activate at boot
<code>disable</code> <i>unit</i>	Prevents <i>unit</i> from activating at boot
<code>isolate</code> <i>target</i>	Changes operating mode to <i>target</i>
<code>start</code> <i>unit</i>	Activates <i>unit</i> immediately
<code>stop</code> <i>unit</i>	Deactivates <i>unit</i> immediately
<code>restart</code> <i>unit</i>	Restarts (or starts, if not running) <i>unit</i> immediately
<code>status</code> <i>unit</i>	Shows <i>unit</i> 's status and recent log entries
<code>kill</code> <i>pattern</i>	Sends a signal to units matching <i>pattern</i>
<code>reboot</code>	Reboots the computer
<code>daemon-reload</code>	Reloads unit files and systemd configuration

Unit file statuses

State	Meaning
bad	Some kind of problem within systemd ; usually a bad unit file
disabled	Present, but not configured to start autonomously
enabled	Installed and runnable; will start autonomously
indirect	Disabled, but has peers in Also clauses that may be enabled
linked	Unit file available through a symlink
masked	Banished from the systemd world from a logical perspective
static	Depended upon by another unit; has no install requirements

Journald and journalctl

- Journald is a system paired with systemd that logs for all the units under systemd control
- Use journalctl to read these logs (located at run/)
- Notable flags: -b X allows you to go back to a previous boot (negative number) or a specific boot (boots hash) -u X allows you to look for a specific unit (service name)

```
$ journalctl
```

```
-- Logs begin at Fri 2016-02-26 15:01:25 UTC, end at Fri 2016-02-26
   15:05:16 UTC. --
Feb 26 15:01:25 ubuntu systemd-journal[285]: Runtime journal is using
   4.6M (max allowed 37.0M, t
Feb 26 15:01:25 ubuntu systemd-journal[285]: Runtime journal is using
   4.6M (max allowed 37.0M, t
Feb 26 15:01:25 ubuntu kernel: Initializing cgroup subsys cpuset
Feb 26 15:01:25 ubuntu kernel: Initializing cgroup subsys cpu
Feb 26 15:01:25 ubuntu kernel: Linux version 3.19.0-43-generic (build@
   lcy01-02) (gcc version 4.
Feb 26 15:01:25 ubuntu kernel: Command line: BOOT_IMAGE=/boot/vmlinuz-
   3.19.0-43-generic root=UUID
Feb 26 15:01:25 ubuntu kernel: KERNEL supported cpus:
Feb 26 15:01:25 ubuntu kernel:   Intel GenuineIntel
...
```

What to do if a system fails to boot?

- Don't debug; just restore the system to a known-good state
 - Always possible assuming you have backups of your system
 - Depending on what you backup and what you need to restore, may be costly
- Bring the system up just enough to run a shell, and debug interactively
 - Use single-user, rescue, or emergency mode to bring up all little as possible before a shell
 - May not work in the case that the bug is very early in the boot
 - Doesn't boot network, so requires physical access to the system
- Boot a separate system image, mount the sick system's filesystem, and investigate
 - Requires the installation of a new system image in addition to the sick one
 - Gives you an external view of the system, which may make some issues more difficult to trace

Shutdown Procedures

- Shutting down a system is governed by three primary commands:
 - `halt`: performs necessary shutdown procedures, logs the shutdown, kills non essential processes, flushes cached filesystem blocks to disk (add `-p` flag to power off the system as well)
 - `reboot`: this performs a halt and then causes the system to boot again
 - `shutdown`: allows for scheduled halts and reboots and prints warnings to users that the system will soon become unavailable

Questions?

Additional Resources

[Digital Ocean: Understanding Systemd Units and Unit Files](#)