


# CS183

Instructor: Ali Davanian

Most slides were adopted from Romit Roy Choudhury and Ishani Janveja (from UC Berkeley)



## L4 Networks and Sockets

# What to learn from Transport Layer?

- Understand principles behind transport layer services:
  - Multiplexing/demultiplexing, reliable data transfer, flow control, congestion control
- Learn about transport layer protocols in the Internet:
  - UDP: connectionless transport
  - TCP: connection-oriented transport
- Learn Socket Programming in Linux
  - Sometimes troubleshooting network services require deep understanding of how sockets work

# Transport vs. network layer

- *Network layer*: logical communication between hosts
- *Transport layer*: logical communication between processes
  - relies on, enhances, network layer services


## Household analogy:

*12 kids sending letters to 12 kids*

- processes = kids
- app messages = letters in envelopes
- hosts = houses
- transport protocol = Ann to Bill
- network-layer protocol = postal service

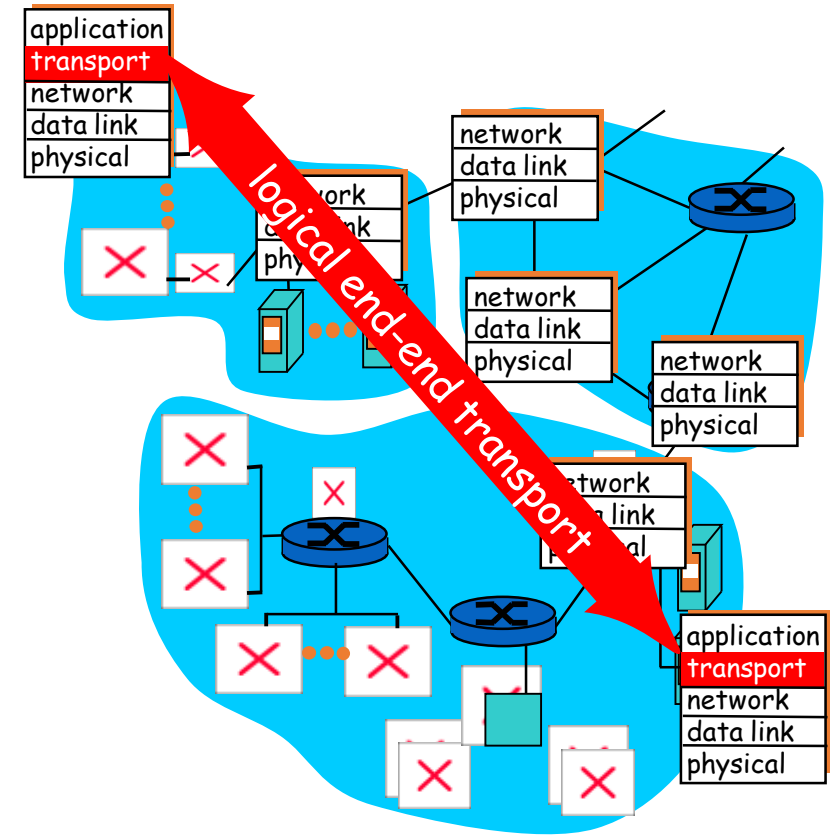
# Transport services and protocols

- Provide logical communication between app processes running on different hosts
- Transport protocols run in end systems
  - Sender: breaks app messages into segments, passes to network layer
  - receiver: reassembles segments into messages, passes to app layer
- More than one transport protocol available to apps
  - Internet: TCP and UDP

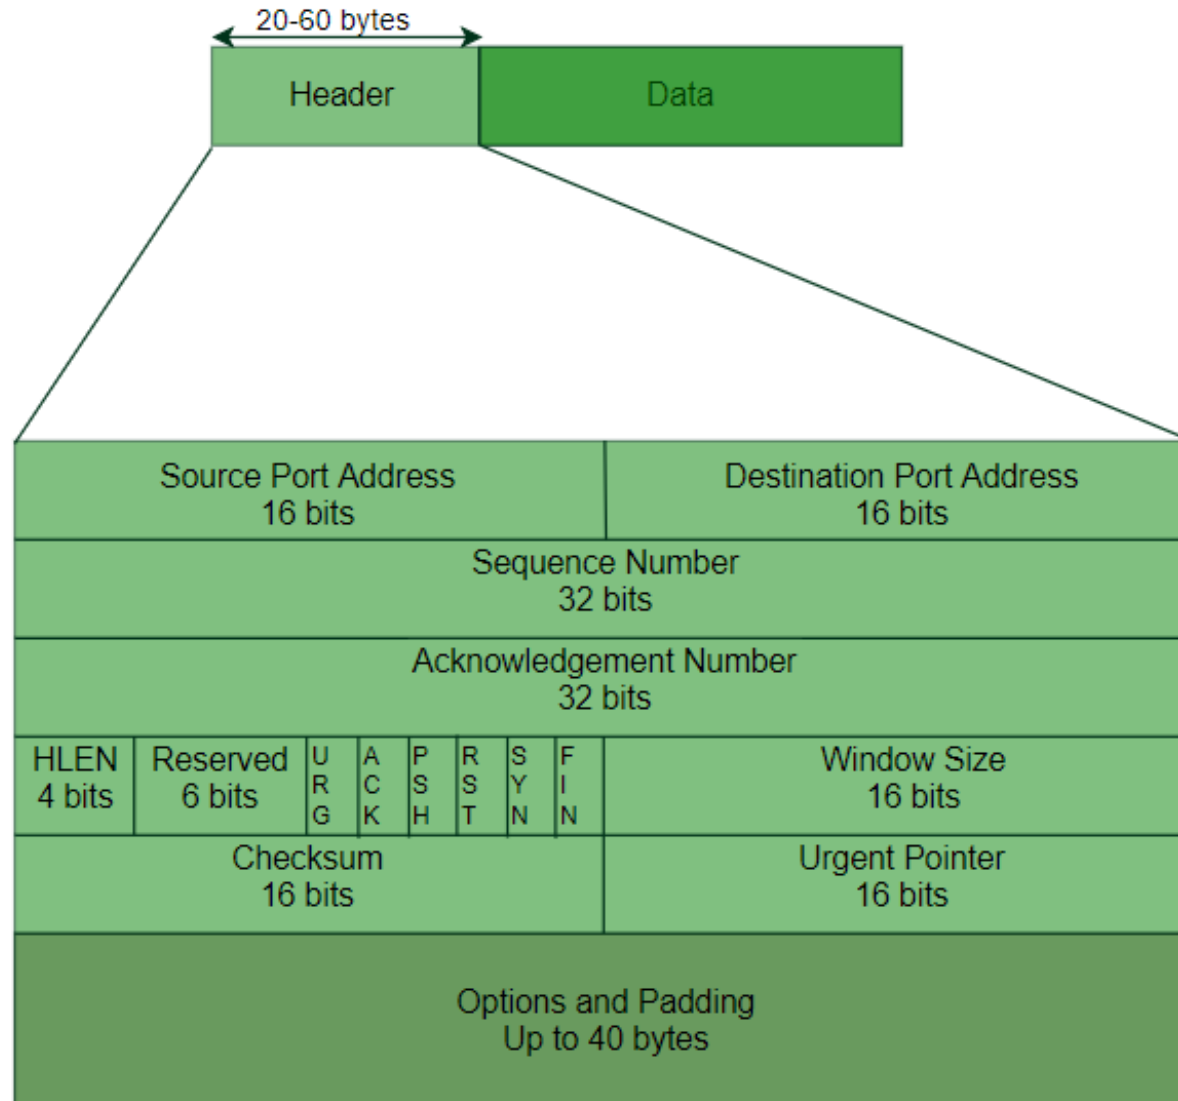


# Transport services and protocols

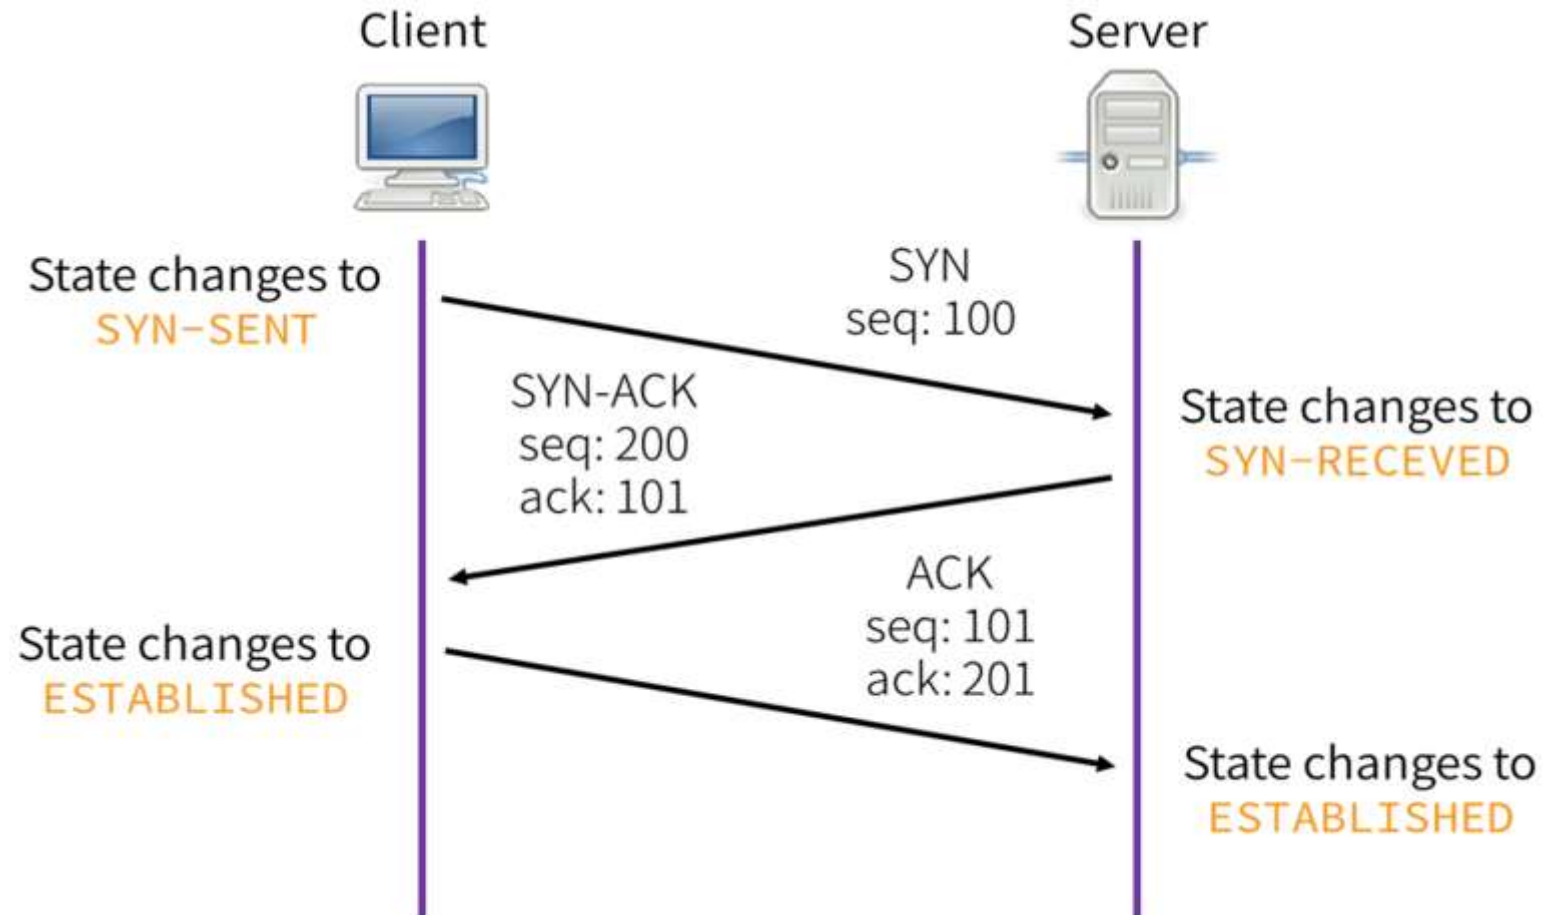
- Reliable, in-order delivery (TCP)
  - congestion control
  - flow control
  - connection setup
- Unreliable, unordered delivery: UDP
  - no-frills extension of “best-effort” IP
- services not available:
  - delay guarantees
  - bandwidth guarantees



# TCP Transport Protocol



# TCP Connection Establishment



# Socket programming

- **Goal:**

- Learn how to build client/server application that communicate using sockets
- Learn what happens in your server/client programs under the hood

## socket

a *host-local*,  
*application-created, OS-controlled*  
interface (a "door") into which  
application process can  
*both send and receive*  
messages to/from another  
application process


## Socket API

- introduced in BSD4.1 UNIX, 1981
- explicitly created, used, released by apps
- client/server paradigm
- two types of transport service via socket API:
  - unreliable datagram
  - reliable, byte stream-oriented



# Socket-programming using TCP

- **Socket:** a door between application process and end-end-transport protocol (UCP or TCP)
- **TCP service:** reliable transfer of bytes from one process to another



# Socket programming: Terminology

Term	Format	Example	Analogy-phone
IP(v4)	x.x.x.x	192.168.0.1	Phone number
Port	N<65535	80	Extension #
Socket	[memory/resource]	s=socket(...)	Physical phone

# Ports, Services, and Hosts

- **Services** running on a **host** bind to a **port**, and use that port as a means of data transmission, which looks like slightly fancy file writing in Linux
- Ports allow us to run many services off the same IP, so 192.168.1.5:80 would handle all the HTTP traffic for the host at 192.168.1.5
- Well known (and permissions regulated) ports are  $0 < x \leq 1024$
- The rest of the port range ( $1025 < x \leq 65535$ ) is open and available
- Here are a few well known ports you should be familiar with

80	443	22	25	53	389	67
HTTP	SSL/TLS	SSH/SFTP	SMTP	DNS	LDAP	DHCP

# Client/server socket interaction: TCP

**Server** (stand-by, waiting for requests)    **Client** (initiate the request)

Create socket  
(Claim resources/ available phone)

Create socket  
(Claim resources/ available phone)

Bind port  
(Claim ID on this machine/  
get a phone extension No.)

-  
(Don't care about source port/ phone #)

**Listen & accept**  
(Listen: Wait for connections/ Wait for phone call  
Accept: Accept connection/ Answer phone call )

**Connect**  
(connect to server IP:PORT/ call phone #))

**Send/ receive**  
(Communication/ Chat on phone)

**Send/ receive**  
(Communication/ Chat on phone)

Close socket  
(End communication/ Hang up the phone)

Close socket  
(End communication/ Hang up the phone)

**Red words: wait for the other side**

# Socket programming with TCP

- **Client must contact server**

- server process must first be running
- server must have created socket (door) that welcomes client's contact

- **Client contacts server by:**

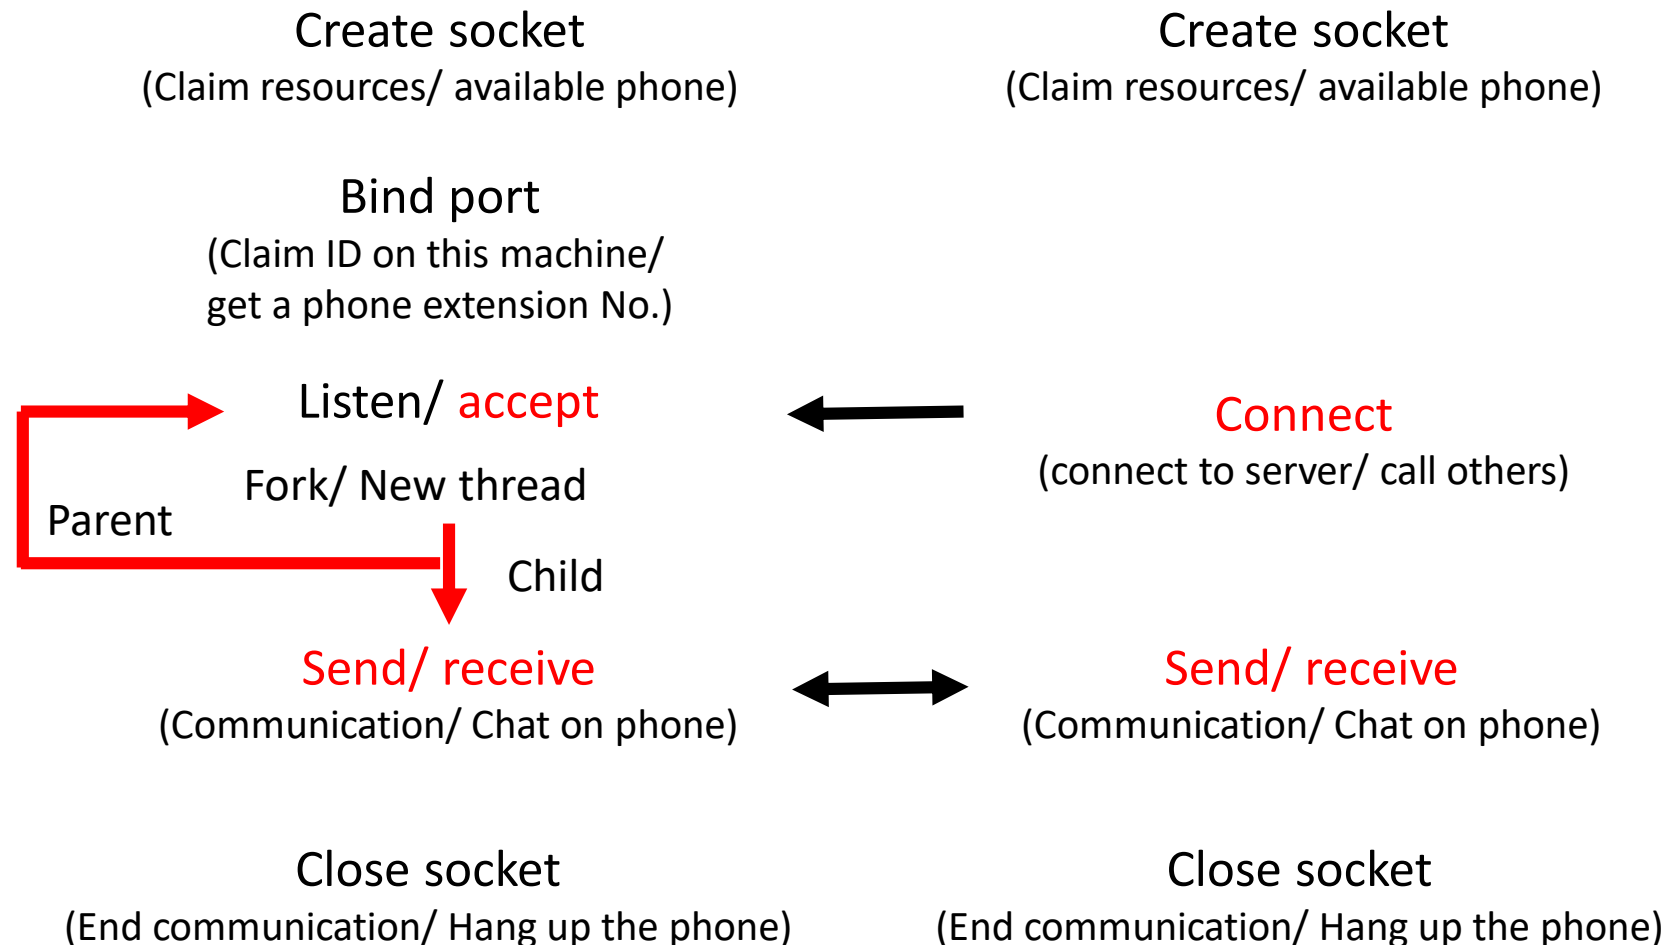
- creating client-local TCP socket
- specifying IP address, port number of server process
- When client creates socket: client TCP establishes connection to server TCP

❖ When contacted by client, **server TCP creates new socket** for server process to communicate with client

- allows server to talk with multiple clients
- source port numbers used to distinguish clients

# Client/server socket interaction: TCP

**Server** (stand-by, waiting for requests)    **Client** (initiate the request)




# Multiplexing/Demultiplexing

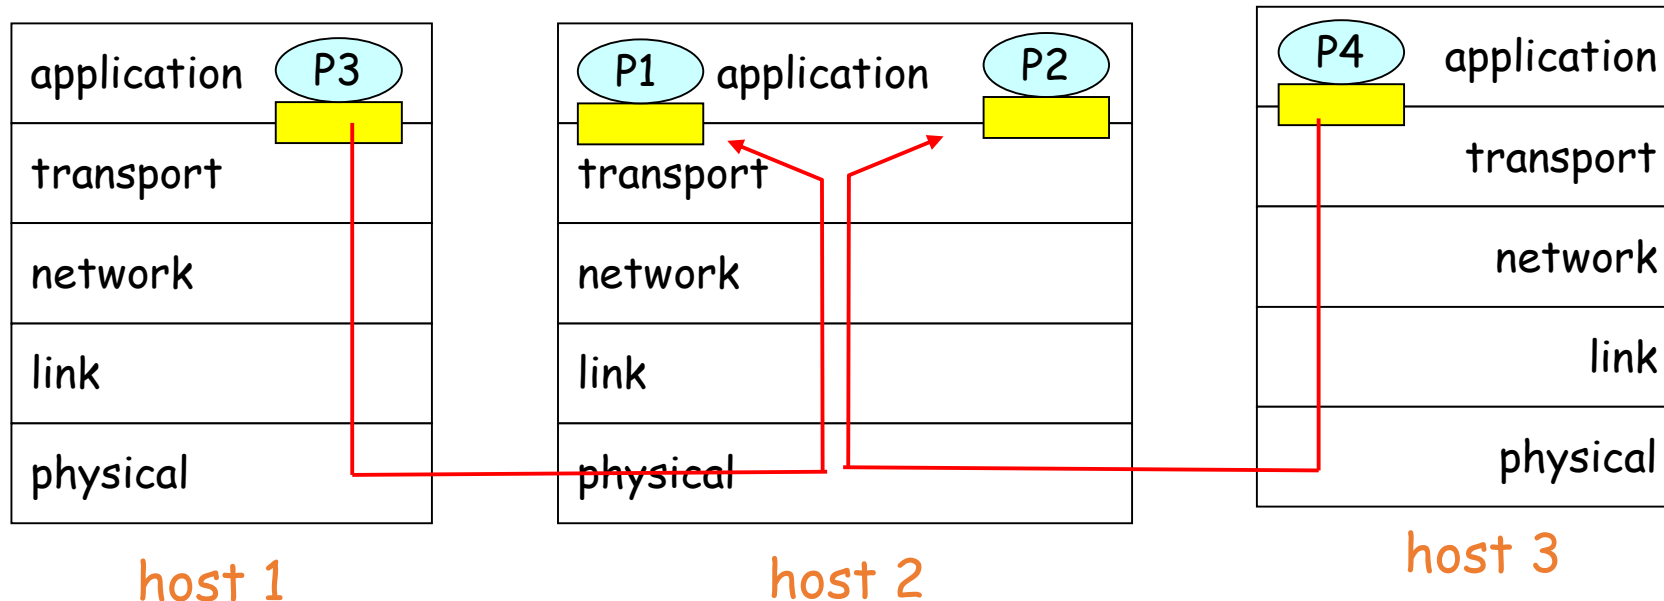
## Demultiplexing at rcv host:

delivering received segments  
to correct socket

## Multiplexing at send host:

gathering data from multiple  
sockets, enveloping data with  
header (later used for  
demultiplexing)

 = socket       = process



# Connection-oriented demux

- TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- recv host uses all four values to direct segment to appropriate socket
- Server host may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
- Web servers have different sockets for each connecting client
  - There are exceptions though



# Socket programming Example: Server side

Create socket  
(Claim resources/ available phone)

Bind port  
(Claim ID on this machine/  
get a phone extension No.)

Listen/ **accept**



**Send/ receive**  
(Communication/ Chat on phone)

Close socket  
(End communication/ Hang up the phone)

```
1  #!/usr/bin/env python2
2  # Echo server program
3  import socket
4  import sys
5
6  HOST = ""           # Symbolic name meaning all available interfaces
7  PORT = 50007        # Arbitrary non-privileged port
8  try:
9      s=socket.socket(socket.AF_INET, socket.SOCK_STREAM)
10 except OSError as msg:
11     s = None
12 try:
13     s.bind((HOST, PORT))
14     s.listen(128)
15 except OSError as msg:
16     s.close()
17     s = None
18 if s is None:
19     print('could not open socket')
20     sys.exit(1)
21 try:
22     conn, addr = s.accept()
23     print('Connected by', addr)
24     while True:
25         data = conn.recv(1024)
26         if not data:
27             break
28         conn.send(data)
29 except:
30     print("Exception")
```

# Testing the example

- Allow connections to the port
  - Depends on the previous configurations (the chains etc.)
  - Fast solution (Dangerous, don't do this in production environments):
    - `sudo iptables -F`
- Run the script:
  - `python server.py`
- Check the port is listened:
  - `sudo ss -tulpn | grep LISTEN`
- Connect to the port:
  - `nc [IP] 50007`

# Socket programming Example: Client side

```
1  #!/usr/bin/env python2
2  #Echo client program
3  import socket
4  import sys
5
6  HOST = '192.168.1.14'    # The remote host
7  PORT = 50007             # The same port as used by the server
8  s = None
9  try:
10     s=socket.socket(socket.AF_INET, socket.SOCK_STREAM)
11 except OSError as msg:
12     s = None
13 try:
14     s.connect((HOST, PORT))
15 except OSError as msg:
16     s.close()
17     s = None
18 if s is None:
19     print('could not open socket')
20     sys.exit(1)
21 try:
22     s.sendall(b'Hello, world')
23     data = s.recv(1024)
24     print('Received', repr(data))
25 except:
26     s.close()
27     print("Exception")
```

Create socket  
(Claim resources/ available phone)

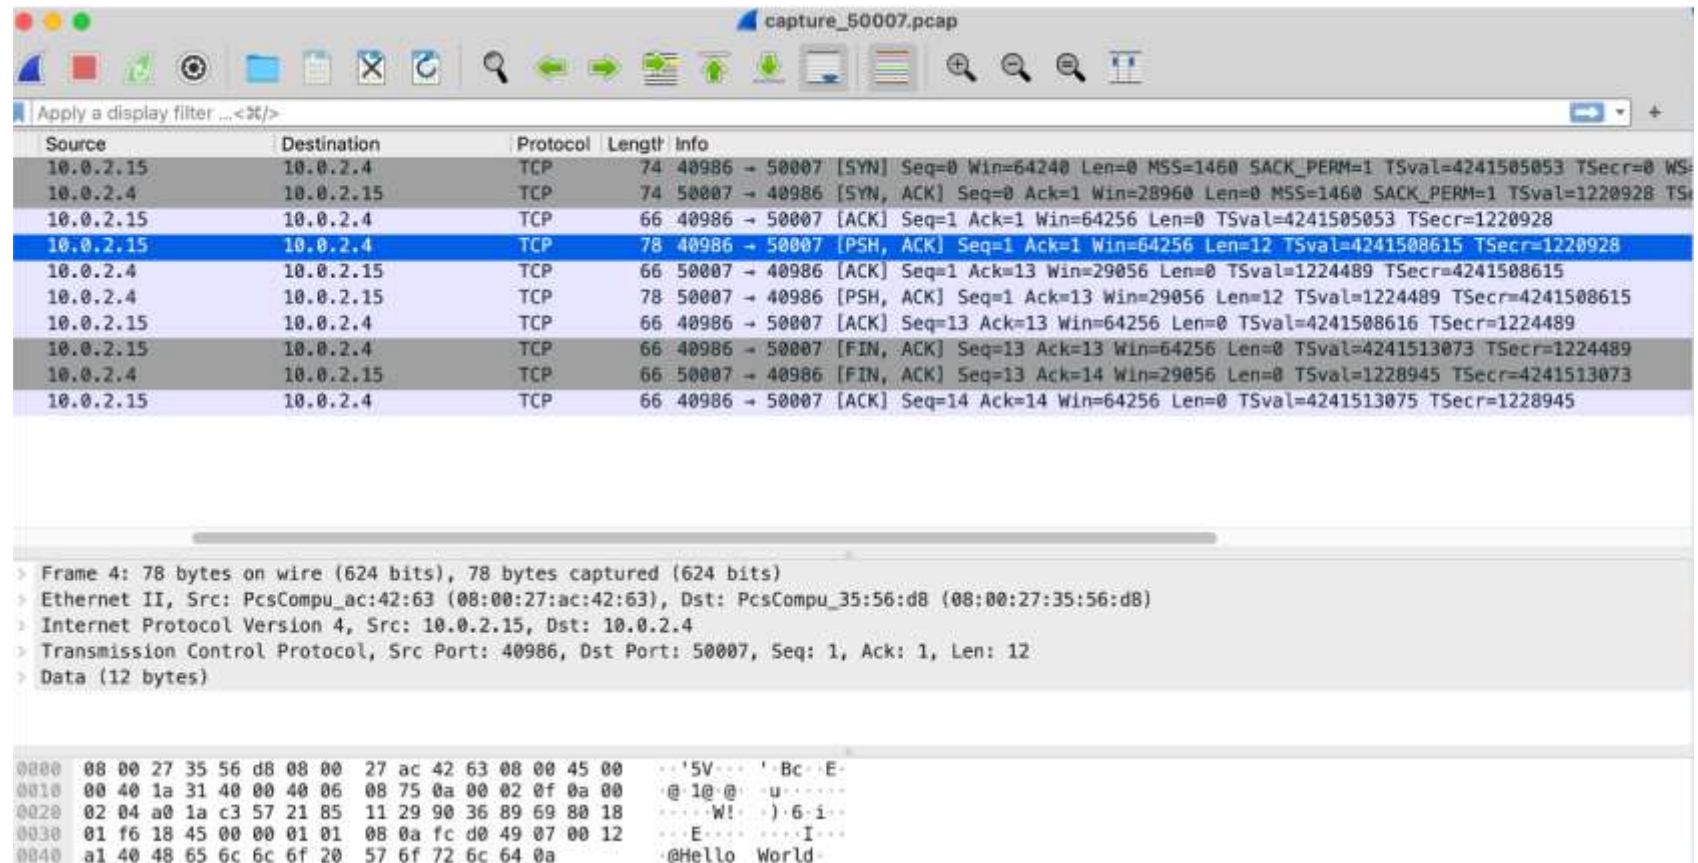
Connect  
(connect to server/ call others)

Send/ receive  
(Communication/ Chat on phone)

Close socket  
(End communication/ Hang up the phone)

# Traffic Analysis

- The communicated traffic can be captured using tcpdump:
  - `sudo tcpdump tcp port 50007 -w capture_50007.pcap`



Source	Destination	Protocol	Length	Info
10.0.2.15	10.0.2.4	TCP	74	40986 → 50007 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 SACK_PERM=1 TSval=4241505053 TSecr=0 WS=
10.0.2.4	10.0.2.15	TCP	74	50007 → 40986 [SYN, ACK] Seq=0 Ack=1 Win=28960 Len=0 MSS=1460 SACK_PERM=1 TSval=1220928 TS
10.0.2.15	10.0.2.4	TCP	66	40986 → 50007 [ACK] Seq=1 Ack=1 Win=64256 Len=0 TSval=4241505053 TSecr=1220928
10.0.2.15	10.0.2.4	TCP	78	40986 → 50007 [PSH, ACK] Seq=1 Ack=1 Win=64256 Len=12 TSval=4241508615 TSecr=1220928
10.0.2.4	10.0.2.15	TCP	66	50007 → 40986 [ACK] Seq=1 Ack=13 Win=29056 Len=0 TSval=1224489 TSecr=4241508615
10.0.2.4	10.0.2.15	TCP	78	50007 → 40986 [PSH, ACK] Seq=1 Ack=13 Win=29056 Len=12 TSval=1224489 TSecr=4241508615
10.0.2.15	10.0.2.4	TCP	66	40986 → 50007 [ACK] Seq=13 Ack=13 Win=64256 Len=0 TSval=4241508616 TSecr=1224489
10.0.2.15	10.0.2.4	TCP	66	40986 → 50007 [FIN, ACK] Seq=13 Ack=13 Win=64256 Len=0 TSval=4241513073 TSecr=1224489
10.0.2.4	10.0.2.15	TCP	66	50007 → 40986 [FIN, ACK] Seq=13 Ack=14 Win=29056 Len=0 TSval=1228945 TSecr=4241513073
10.0.2.15	10.0.2.4	TCP	66	40986 → 50007 [ACK] Seq=14 Ack=14 Win=64256 Len=0 TSval=4241513075 TSecr=1228945

> Frame 4: 78 bytes on wire (624 bits), 78 bytes captured (624 bits)

> Ethernet II, Src: PcsCompu\_ac:42:63 (08:00:27:ac:42:63), Dst: PcsCompu\_35:56:d8 (08:00:27:35:56:d8)

> Internet Protocol Version 4, Src: 10.0.2.15, Dst: 10.0.2.4

> Transmission Control Protocol, Src Port: 40986, Dst Port: 50007, Seq: 1, Ack: 1, Len: 12

> Data (12 bytes)

```
0000 08 00 27 35 56 d8 08 00 27 ac 42 63 08 00 45 00  ..'5V... 'Bc..E-
0010 00 40 1a 31 40 00 40 06 08 75 0a 00 02 0f 0a 00  @1@...u.....
0020 02 04 a0 1a c3 57 21 85 11 29 90 36 89 69 80 18  ....W!...).6-i..
0030 01 f6 18 45 00 00 01 01 08 0a fc d0 49 07 00 12  ...E.....I...
0040 a1 40 48 65 6c 6c 6f 20 57 6f 72 6c 64 0a      @Hello World
```